



# Chapter 4: Intermediate SQL

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Chapter 4: Intermediate SQL

- Joined Relation
  - SQL Data Types and Schemas
  - Integrity Constraints
  - Views
  - Indexes
- 
- Transactions
  - Authorization



# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- Join operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>	<i>Join Conditions</i>
inner join left outer join right outer join full outer join	natural on <predicate> using $(A_1, A_1, \dots, A_n)$



# Join operations – Example

## ■ Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

## ■ Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

## ■ Observe that

prereq information is missing for CS-315 and  
course information is missing for CS-437



# Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.



# Left Outer Join

■ *course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>



# Right Outer Join

■ *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



# Full Outer Join

■ *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101





# Joined Relations – Examples

- **course inner join prereq on**  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above, and a natural join?
- **course left outer join prereq on**  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>



# Joined Relations – Examples

- **course natural right outer join prereq**

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

- **course full outer join prereq using (course\_id)**

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



# Built-in Data Types in SQL

- **date**: Dates, containing a (4 digit) year, month and date
  - Example: **date** '2005-7-27'
- **time**: Time of day, in hours, minutes and seconds.
  - Example: **time** '09:00:30'      **time** '09:00:30.75'
- **timestamp**: date plus time of day
  - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval**: period of time
  - Example: interval '1' day
  - Subtracting a date/time/timestamp value from another gives an interval value
  - Interval values can be added to date/time/timestamp values
- **date, time functions**:
  - `current_date()`, `current_time()`
  - `year(x)`, `month(x)`, `day(x)`, `hour(x)`, `minute(x)`, `second(x)`



# User-Defined Types

- **create type** construct in SQL creates user-defined type

**create type** *Dollars* as numeric (12,2) final

- **create table** *department*  
(*dept\_name* **varchar** (20),  
*building* **varchar** (15),  
*budget* *Dollars*);



# Domains

- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.
- **create domain** *degree\_level* **varchar**(10)  
**constraint** *degree\_level\_test*  
**check** (**value in** ('Bachelors', 'Masters', 'Doctorate'));



# Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
  - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
  - **clob**: character large object -- object is a large collection of character data
  - When a query returns a large object, a pointer is returned rather than the large object itself.



# Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
  - A checking account must have a balance greater than \$10,000.00
  - A salary of a bank employee must be at least \$4.00 an hour
  - A customer must have a (non-null) phone number



# Integrity Constraints on a Single Relation

- not null
- primary key
- unique
- **check** (P), where P is a predicate
- foreign key





# Not Null and Unique Constraints

## ■ not null

- Declare *name* and *budget* to be **not null**

*name* **varchar(20) not null**

*budget* **numeric(12,2) not null**

## ■ unique ( $A_1, A_2, \dots, A_m$ )

- The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a super key (  $\times$  candidate key ) .
- Candidate keys are permitted to be null (in contrast to primary keys).



# The check clause

## ■ **check** (P)

where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```



# Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
  - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”.
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.



# Cascading Actions in Referential Integrity

- **create table** *course* (  
    *course\_id* **char**(5) **primary key**,  
    *title* **varchar**(20),  
    *dept\_name* **varchar**(20) **references** *department*  
)
- **create table** *course* (  
    ...  
    *dept\_name* **varchar**(20),  
    **foreign key** (*dept\_name*) **references** *department*  
        **on delete cascade**  
        **on update cascade**,  
    ...  
)
- alternative actions to cascade: **set null**, **set default**, **restricted**

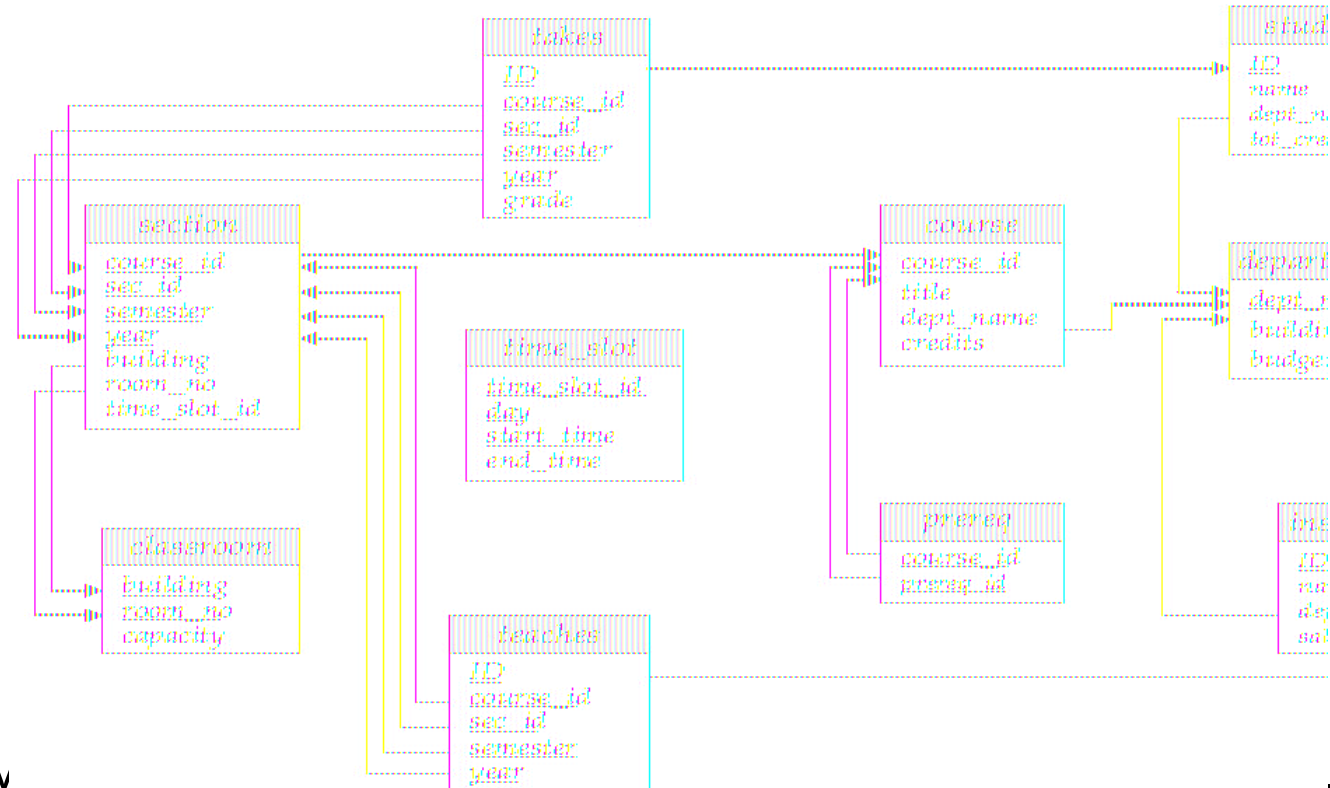


# Integrity Constraint Violation During Transactions

- **create table** *person* (  
    *ID* **char**(10),  
    *name* **char**(40),  
    *mother* **char**(10),  
    *father* **char**(10),  
    **primary key** (*ID*),  
    **foreign key** (*father*) **references** *person*,  
    **foreign key** (*mother*) **references** *person*);
  
- How to insert a tuple without causing constraint violation ?
  - insert father and mother of a person before inserting person
  - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
  - OR defer constraint checking to **transaction end**.



- **check** ((*course\_id*, *sec\_id*, *semester*, *year*) in  
          (select *course\_id*, *sec\_id*, *semester*, *year* from teaches))





# Complex Check Clauses

- Unfortunately: subquery in check clause not supported by pretty much any database
  - Alternative: **triggers**



# assertion

- **create assertion** <assertion-name> **check** <predicate>;
- **create assertion** *credits\_earned\_constraint* **check**  
(not exists  
    (select *ID*  
    from *student*  
    where *tot\_cred* <> (  
        select sum(*credits*)  
        from *takes* natural join *course*  
        where *student.ID=takes.ID*  
            and *grade* is not null  
            and *grade*<>'F'))





# Views

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Consider a person who needs to know **an instructors name and department, but not the salary**. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.



# View Definition

- A view is defined using the **create view** statement which has the form

**create view** *v* **as** < query expression >

where <query expression> is any legal SQL expression. The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.



# Example Views

- A view of instructors without their salary

```
create view faculty as  
  select ID, name, dept_name  
  from instructor
```

- Find the names of all instructors in the Biology department

```
select name  
from faculty  
where dept_name = 'Biology'
```

- Create a view of department salary totals

```
create view departments_total_salary(dept_name, total_salary) as  
  select dept_name, sum (salary)  
  from instructor  
  group by dept_name;
```



# Views Defined Using Other Views

- **create view** *physics\_fall\_2009* **as**  
    **select** *course.course\_id, sec\_id, building, room\_number*  
    **from** *course, section*  
    **where** *course.course\_id = section.course\_id*  
          **and** *course.dept\_name = 'Physics'*  
          **and** *section.semester = 'Fall'*  
          **and** *section.year = '2009';*
- **create view** *physics\_fall\_2009\_watson* **as**  
    **select** *course\_id, room\_number*  
    **from** *physics\_fall\_2009*  
    **where** *building= 'Watson';*



# View Expansion

- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as  
(select course_id, room_number  
from (select course.course_id, building, room_number  
      from course, section  
      where course.course_id = section.course_id  
           and course.dept_name = 'Physics'  
           and section.semester = 'Fall'  
           and section.year = '2009')  
where building = 'Watson';
```



# Update of a View

- Add a new tuple to *faculty* view which we defined earlier

```
create view faculty as  
    select ID, name, dept_name  
    from instructor
```

- **insert into** *faculty* **values** ('30765', 'Green', 'Music');

This insertion must be represented by the insertion of the tuple

('30765', 'Green', 'Music', null)

into the *instructor* relation



# Some Updates cannot be Translated Uniquely

- **create view** *instructor\_info* as  
    **select** *ID, name, building*  
    **from** *instructor, department*  
    **where** *instructor.dept\_name= department.dept\_name;*
- **insert into** *instructor\_info* **values** ('69987', 'White', 'Taylor');
  - ▶ which department, if multiple departments in Taylor?
- Most SQL implementations allow updates only on simple views(**updatable views**)
  - The **from** clause has only one database relation.
  - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
  - Any attribute not listed in the **select** clause can be set to null
  - The query does not have a **group** by or **having** clause.



## \* Materialized Views

- **Materializing a view**: create a physical table containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the materialized view result becomes out of date
  - Need to **maintain the view**, by updating the view whenever the underlying relations are updated.





# \* View and Logical Data Independence

- If relation **S(a, b, c)** is split into two sub relations **S1(a,b)** and **S2(a,c)**  
How to realize the logical data independence?

- 1) create table S1 ...;  
create table S2 ...;
- 2) insert into S1 select a, b from S;  
insert into S2 select a, c from S;
- 3) drop table S;
- 4) create view S(a,b,c) as select a,b,c from S1 natural join S2;

select \* from S where ... → select \* from S1 natural join S2 where ...

insert into S values (1 ,2,3) → insert into S1 values (1, 2);  
insert into S2 values (1 ,3);



# Indexes

- **create table** *student*  
(     *ID* **varchar** (5),  
      *name* **varchar** (20) **not null**,  
      *dept\_name* **varchar** (20),  
      *tot\_cred* **numeric** (3,0) **default** 0,  
      **primary key** (*ID*) )
  - **create index** *studentID\_index* **on** *student*(*ID*)
  - Indices are data structures used to speed up access to records with specified values for index attributes
    - e.g. **select** \*  
      **from** *student*  
      **where** *ID* = '12345'
- can be executed by using the index to find the required record, without looking at all records of *student*
- More on indices in Chapter 11*



# Transactions

- Unit of work (**NONE** or **ALL**)
- Atomic transaction
  - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
  - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
  - Can turn off auto commit for a session (e.g. using API)
  - In MySQL:
    - >**SET AUTOCOMMIT=0;**
  - In SQL:1999, can use: **begin atomic .... end**
    - ▶ Not supported on most databases



# Transactions

- Transaction example :
- **SET AUTOCOMMIT=0**

```
UPDATE account SET balance=balance-100 WHERE ano='1001';  
UPDATE account SET balance=balance+100 WHERE ano='1002';  
COMMIT;
```

```
UPDATE account SET balance=balance -200 WHERE ano='1003';  
UPDATE account SET balance=balance+200 WHERE ano='1004';  
COMMIT;
```

```
UPDATE account SET balance=balance+balance*2.5%;  
COMMIT;
```



# ACID Properties

- A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:
  - **Atomicity**. Either all operations of the transaction are properly reflected in the database or none are.
  - **Consistency**. Execution of a transaction in isolation preserves the consistency of the database.
  - **Isolation**. Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
    - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.
  - **Durability**. After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



# Authorization( 授权 )

Forms of authorization on parts of the database:

- **Select** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.

Forms of authorization to modify the database schema

- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.



# Authorization Specification in SQL

- The **grant** statement is used to confer authorization
  - grant** <privilege list> // **privilege** : 权限
  - on** <relation name or view name> **to** <user list>
- <user list> is:
  - a user-id
  - **public**, which allows all valid users the privilege granted
  - A role (more on this later)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).



# Privileges in SQL

- **grant select on *instructor* to  $U_1, U_2, U_3$**
- **grant select on *department* to public**
- **grant update (*budget*) on *department* to  $U_1, U_2$**
- **grant all privileges on *department* to  $U_1$**





# Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.  
**revoke** <privilege list>  
**on** <relation name or view name>  
**from** <user list>
- Example:  
**revoke select on *branch* from  $U_1, U_2, U_3$**
- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.
- All privileges that depend on the privilege being revoked are also revoked.



# Roles

- **create role** instructor;
- **grant** *instructor* **to** **Amit**;
- Privileges can be granted to roles:
  - **grant select on** *takes* **to** *instructor*;
- Roles can be granted to users, as well as to other roles
  - **create role** *teaching\_assistant*
  - **grant** *teaching\_assistant* **to** *instructor*;
    - ▶ *Instructor* inherits all privileges of *teaching\_assistant*
- Chain of roles
  - **create role** *dean*;
  - **grant** *instructor* **to** *dean*;
  - **grant** *dean* **to** Satoshi;



# Authorization on Views

- **create view** *geo\_instructor* **as**  
(**select** \*  
  **from** *instructor*  
  **where** *dept\_name* = 'Geology');
- **grant select on** *geo\_instructor* **to** *geo\_staff*



# Other Authorization Features

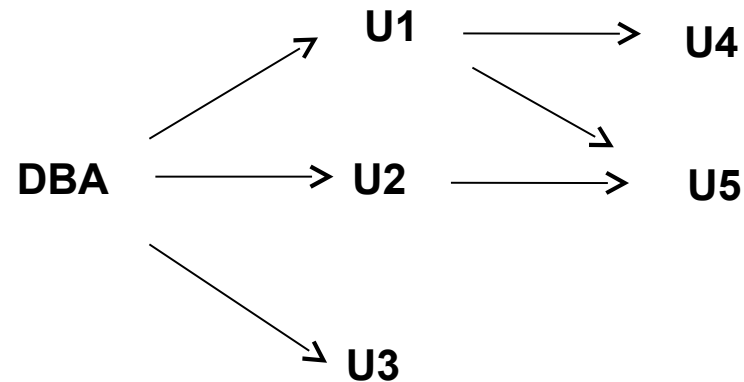
- **references** privilege to create foreign key
  - **grant reference** (*dept\_name*) **on** *department* **to** Mariano;
  - why is this required?





# Other Authorization Features

- transfer of privileges
  - **grant select on *department* to Amit with grant option;**
  - **revoke select on *department* from Amit, Satoshi cascade;**
  - **revoke select on *department* from Amit, Satoshi restrict;**





# End of Chapter 4

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use